

Class10: Halloween Mini-Project

Shazreh Hassan (PID: A13743949)

Importing candy data

```
candy_file <- read.csv("https://raw.githubusercontent.com/fivethirtyeight/data/master/candy-power-ranking.csv")
candy <- read.csv("https://raw.githubusercontent.com/fivethirtyeight/data/master/candy-power-ranking.csv")
head(candy)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crispedrice	wafer
100 Grand	1	0	1			0	0	1
3 Musketeers	1	0	0			0	1	0
One dime	0	0	0			0	0	0
One quarter	0	0	0			0	0	0
Air Heads	0	1	0			0	0	0
Almond Joy	1	0	0			1	0	0
	hard	bar	pluribus	sugarpercent	pricepercent	pricepercent	winpercent	
100 Grand	0	1	0	0.732	0.860	0.860	66.97173	
3 Musketeers	0	1	0	0.604	0.511	0.511	67.60294	
One dime	0	0	0	0.011	0.116	0.116	32.26109	
One quarter	0	0	0	0.011	0.511	0.511	46.11650	
Air Heads	0	0	0	0.906	0.511	0.511	52.34146	
Almond Joy	0	1	0	0.465	0.767	0.767	50.34755	

```
flextable:: flextable(head(candy))
```

chocolate	fruity	caramel	peanut	almond	nougat	crispedrice	wafer	hard	bar	pluribus	sugarpercent
1	0	1	0	0	0	1	0	0	1	0	66.97173
1	0	0	0	0	1	0	0	0	1	0	67.60294

chocolate	fruity	caramel	peanut	yalmond	nougat	crisped	rice	wafer	hard	bar	pluribus	s
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	0	0	0	0	0	0
1	0	0	1	0	0	0	0	0	0	1	0	0

Q1. How many different candy types are in this dataset?

```
library(dplyr)
nrow(candy)
```

[1] 85

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

[1] 38

What is your favorite candy?

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

```
#favorite candy: Junior Mints
candy["Junior Mints", ]$winpercent
```

[1] 57.21925

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat", ]$winpercent
```

[1] 76.7686

Q5. What is the winpercent value for "Tootsie Roll Snack Bars"?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

```
library("skimr")
skim(candy)
```

Table 2: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondyn	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

winpercent looks like it is on a 1-100 scale while the others are 0-1


Q7. What do you think a zero and one represent for the candy\$chocolate column?

Zero means the candy does not have chocolate, and 1 means it does.

Q8. Plot a histogram of winpercent values

```
library(ggplot2)


ggplot(candy) +
  aes(winpercent) +
  geom_histogram(bins=20)
```



Q9. Is the distribution of winpercent values symmetrical?

```
#density plot

ggplot(candy) +
  aes(winpercent) +
  geom_density()
```



No, it looks skewed to the left of the 50% mark.

Q10. Is the center of the distribution above or below 50%?

```
summary(candy$winpercent)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	22.45	39.14	47.83	50.32	59.86	84.18

The center of the distribution is below 50%.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
# 1. find all chocolate candy in the dataset
ind.choc <- candy$chocolate==1
choc.candy <- candy[ind.choc,]

# 2. extract their winpercent values
choc.win <- choc.candy$winpercent

# 3. find the mean of these values
choc.mean <- mean(choc.win)
```

```
# 4-6. do the same for fruity candy
fruit.win <- candy[candy$fruity==1,]$winpercent
fruit.mean <- mean(fruit.win)

# 7. which mean value is higher?
choc.mean
```

[1] 60.92153

```
fruit.mean
```

[1] 44.11974

Chocolate is higher ranked

Q12. Is this difference statistically significant?

```
t.test(choc.win, fruit.win)
```

```
Welch Two Sample t-test

data: choc.win and fruit.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

Yes, this difference is statistically significant.

Overall candy rankings

Q13. What are the five least liked candy types in this set?

```
ord.ind <- order(candy$winpercent)
head(candy[ord.ind,], 5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat				
Nik L Nip	0	1	0	0	0	0				
Boston Baked Beans	0	0	0	1	0	0				
Chiclets	0	1	0	0	0	0				
Super Bubble	0	1	0	0	0	0				
Jawbusters	0	1	0	0	0	0				
	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip	0	0	0	1			0.197		0.976	
Boston Baked Beans	0	0	0	1			0.313		0.511	
Chiclets	0	0	0	1			0.046		0.325	
Super Bubble	0	0	0	0			0.162		0.116	
Jawbusters	0	1	0	1			0.093		0.511	
	win	percent								
Nik L Nip	22.44534									
Boston Baked Beans	23.41782									
Chiclets	24.52499									
Super Bubble	27.30386									
Jawbusters	28.12744									


Q14. What are the top 5 all time favorite candy types out of this set?

```
tail(candy[ord.ind,], 5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat		
Snickers	1	0	1	1	1	1		
Kit Kat	1	0	0	0	0	0		
Twix	1	0	1	0	0	0		
Reese's Miniatures	1	0	0	1	0	0		
Reese's Peanut Butter cup	1	0	0	1	0	0		
	crisped	rice	wafer	hard	bar	pluribus	sugar	percent
Snickers	0	0	1	0	0	0	0.546	
Kit Kat	1	0	1	0	0	0	0.313	
Twix	1	0	1	0	0	0	0.546	
Reese's Miniatures	0	0	0	0	0	0	0.034	
Reese's Peanut Butter cup	0	0	0	0	0	0	0.720	
	price	percent	win	percent				
Snickers	0.651	76.67378						
Kit Kat	0.511	76.76860						
Twix	0.906	81.64291						
Reese's Miniatures	0.279	81.86626						
Reese's Peanut Butter cup	0.651	84.18029						


Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy)+  
  aes(winpercent, reorder(rownames(candy),winpercent))+  
  geom_col()
```



Add some color

```
mycols <- rep("black", nrow(candy))  
mycols[candy$chocolate==1] <- "chocolate"  
mycols[candy$bar==1] <- "brown"  
mycols[candy$fruity==1] <- "red"  
  
ggplot(candy)+  
  aes(winpercent, reorder(rownames(candy),winpercent))+  
  geom_col(fill=mycols)
```




Winpercent vs Pricepercent

```
library(ggrepel)

ggplot(candy)+
  aes(winpercent, pricepercent, label=rownames(candy))+
  geom_point(col=mycols)+
  geom_text_repel(col=mycols, size=3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps




Exploring the correlation structure

```
cij <- cor(candy)
```

```
library(corrplot)
```

corrplot 0.95 loaded

```
corrplot(cij)
```



Principal component analysis

The main function in base R for this is `prcomp()` and we want to set `scale=TRUE` here:

```
pca <- prcomp(candy, scale=TRUE)  
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

Let's visualize our first main result figure - the PC plot or PC1 vs PC2

```
ggplot(pca$x) +
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point(col=mycols) +
  geom_text_repel(col=mycols) +
  labs(title="Shaz's candy space map")
```

Warning: ggrepel: 54 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Don't forget about your variable "loadings" - how the original variables contribute to your new PCs...

```
ggplot(pca$rotation) +
  aes(PC1, rownames(pca$rotation)) +
  geom_col()
```

